Recruiter Summary

Candidate highlight As per Current Job Description

M.P Prasad is a senior-level professional with 5+ years of experience, currently in a Sr Gen AI Developer role. Their career is heavily focused on AI/ML development, data science, and building/deploying intelligent applications on major cloud platforms.

Core technical strengths include deep proficiency in Python, modern infrastructure tooling (Docker, Kubernetes, CI/CD), and experience with distributed computing frameworks, demonstrating strong foundational skills for automation and systems work.

Verdict:

■ Low overall match for this specific SRE role due to a fundamental career focus misalignment.
■ Candidate's profile is strongly oriented towards AI/ML development, not Site Reliability Engineering or infrastructure operations.
■ Consider only if seeking a candidate for a hard pivot, but significant gaps in core SRE practices and Oracle Cloud make them a non-ideal fit.

Overall Matching Score: Overall Match: 1.8 / 5 (≈ 36%)

Skill Area	JD Requirement	Resume Evidence	Score	Match Summary
Programming (Python)	Required proficiency in Python.	Explicitly listed as a primary skill.	5/5	Strong match.
Experience & Seniority	10+ years, IC5 level	5+ years of experience shown, currently in a senior developer role.	3/5	Below required tenure, but at senior level.
Location / Visa	Bangalore, India. No visa sponsorship.	Indian phone number provided; location not specified but presumed local.	5/5	Likely aligns with location and visa requirements.
Professional Curiosity	Desire to develop deep understanding (Optional).	Engagement with wide array of cutting-edge AI/ML tech (LLMs, Transformers, etc.).	5/5	Excellent match for this trait.
Automation	Build automation to prevent problem recurrence.	Skills in CI/CD, Docker, Kubernetes, Azure DevOps.	4/5	Strong indirect evidence via DevOps tooling.
Orchestration	Understanding of orchestration principles.	Kubernetes listed as a skill.	4/5	Direct evidence with key tool.
Problem Solving	Solve complex infrastructure problems.	Inherent in complex AI/ML development roles.	4/5	Strong general aptitude evidenced.
Distributed Systems	Design for large-scale distributed systems.	Experience with Sparks-Hadoop and cloud-deployed services.	3/5	Some relevant adjacent experience.

Skill Area	JD Requirement	Resume Evidence	Score	Match Summary
Service Architecture	Guide improvements in service architecture.	Experience building AI services implies some architecture, but focus is different.	3/5	Indirect, conceptual overlap.
Production Troubleshooting	Ultimate escalation for complex production issues.	Senior developer debugging implied, but not operational troubleshooting focus.	2/5	Limited evidence for required depth.
System Performance	Performance analysis, system tuning.	AI/ML model tuning experience, not system/software performance.	2/5	Adjacent but not equivalent skill.
Capacity Planning	Service capacity planning and demand forecasting.	Not explicitly mentioned in profile.	1/5	No direct evidence.
Career Focus & Role Fit	SRE with full-stack service ownership, operational focus.	Career focused on AI/ML development and data science, not SRE or infrastructure operations.	1/5	Major misalignment in core role purpose.
SRE Concepts	Site Reliability Engineering practices.	No direct mention of SRE, SLAs, error budgets, or production support.	1/5	Critical gap; no evidenced experience.
Cloud Platform	Oracle Cloud experience.	Lists AWS, GCP, Azure. No mention of Oracle Cloud.	0/5	No direct evidence; gap.

M.P Prasad

Sr Gen AI Developer (RAG's – GenAI – AgenticAI)

☎: +91 6363185714
: prasad.agenticai.dev@gmail.com
in: linkedin.com/in/prasad-337412278

Summary: Experienced Data Scientist with 5.6+ years **(3.5+ in Gen AI)** in developing and implementing AI-agentic, RAG, LangChain, LangGraph, and LLM-driven applications for NLP tasks. Skilled in integrating vector storage, SQL databases, and cloud platforms like AWS and **Azure**, I built intelligent data-driven solutions also designing, optimizing, and deploying scalable AI workflows, with expertise in end-to-end project delivery that drives automation, efficiency, and impactful business outcomes.

Professional Experience Details:

- NTT Data Info Pvt Ltd || Bangalore || Feb 2025 to (Present) as a **Sr Gen AI developer** (Under NLB Services Pvt Ltd)
- ExamRoom.AI India Pvt Ltd || Bangalore || Sep 2023 to Dec 2024 as a Sr Data Scientist
- NV Data Systems Pvt Ltd || Bangalore || June 2020 to Aug 2023 as a Data Scientist

Technical Stacks:

Programming: Python • NumPy • Pandas • Scikit-learn • NLP • SQL • TensorFlow • Keras • PyTorch • OpenCV • CUDA
Machine Learning: Linear Regression • Logistic Regression • Decision Trees • Random Forest • Naive Bayes • K-NN • K- Means • Clustering Algorithms • Boosting Algorithms • Principle Component Analysis • Ensemble techniques.
Deep Learning: Deep-NN • RNN • CNN • LSTM • Transformers • Sparks-Hadoop • GANs • YOLO • MongoDB •
LLM's & Generative AI: Large Language Models(LLM's) • LLM's-Fine Tune-Testing-Merging • Pre-training • GGUF • PEFT's • LoRA • Q-LoRA • Prompt-Eng • DSPy • Text Generation • Unsloth-Training • Machine Translation • AWS • SQL
Management Tools: Project management (Jira) • GitLab • Microsoft Teams • Docker • Kubernetes • CI-CD • Azure Devops
Generative AI Tools:- AI-Agents, Lengchain, LangGrah, Azure OpenAI & Rag, ChatGpt, ADK, MCP, GCP, Vertex-AI

Education Details:

- Bachelor of Technology in Electrical & Electronics Engineering @ Jawaharlal Nehru Technological University-2012

Certification Courses:

- IBM-Rag-and-Agentic-Ai Professional certificate from Coursera with IBM.
- Large Language Model Operations (LLMOps) Specialization from Coursera with Duke University.

Responsibilities:

- Pre-processed structured/unstructured data from SQL and other sources for LLM, RAG, and agentic AI training.
- Implemented RAG workflows with vector storage solutions Chroma, enabling efficient semantic search and knowledge retrieval. Enhanced contextIal text generation with Azure openAI models to improving accuracy and latency.
- Built and fine-tuned Large Language Models (Mistral, Llama) for NLP tasks like summarization, Question-Answering systems.
- Developed Agentic-AI workflows leveraging LangGraph for orchestrating multi-step reasoning, and intelligent task execution.
- Optimized models into GGUF formats (Q-LoRA, PEFT) for efficient inference, lightweight deployment, and faster execution
- Integrated LLM pipelines with SQL-driven automation applications and enterprise data sources, enabling intelligent query handling, automated reporting, and enhanced decision-making efficiency across business workflows.
- Collaborate with cross-functional teams to deliver scalable AI solutions & alignment with business goals and measurable impact.

PROJECTS:-

Project: Text Gen for Various (4) Types of Question Gen Using LLMs (Mistral-7B) (Nov-23 to Sept-24)

LLM's used:

- Python
- Mistral-7B-instruct
- Finetune & Pretrain
- Langchain, GGUF's

- Developed a sophisticated text generation system to create five types of questions using the Mistral-7B Instruct model & also I used different LLM's for various question model.
- Prepared and collected a sufficient large dataset and performed pre-processing also prepared finetuned Dataset format to ensure high relevance and contextual accuracy.
- Fine-tuned model with the pre-processed data, optimize it for question generation tasks.
- Integrated the fine-tuned model into an application, enabling dynamic question generation based on user input.
- Conducted extensive testing and validation to ensure the model's performance and reliability in various scenarios.

Project: Object Detection & Tracking with YOLO (oct-24 to dec-24)

Model's used:

- Python / C++
- YOLO/ONNX
- Retina-Face
- Resnet / Face-net

- Performed extensive data collection and preprocessing to train the model, ensuring accurate and performed Object detection for 34 suspected classes
- Deployed model on cloud platforms, enabling efficient and scalable translation services.
- Optimized YOLO model using and hyperparameter, achieving improve accuracy & faster inference.
- Implemented object tracking with OpenCV, enhancing detection performance for video streams.

Project: Speech Emotion Recognition (Nov-22 to April-23)

Model used:

- Python
- Wave2vec model
- Finetuned
- Deep learning

- Developed a speech emotion recognition system utilizing the Wav2Vec model to leverage state-of-the-art speech processing techniques. Collected and processed audio data from four diverse datasets to train the model, ensuring various emotional expressions and contexts.
- Fine-tuned the Wav2Vec model on the pre-processed data, achieving high accuracy in detecting and classifying emotions from speech.
- Integrated the trained Wav2Vec model into a customer service application, enhancing real-time emotion recognition and response. This improved user experience & increased customer satisfaction.

Project: Document Processing & RAG Pipeline using Azure Document Intelligence (Feb-25 to June-25)

Tools used:

- Python, Pyodbc
- Azure-OpenAI
- Flask, Bs4, fitz
- SQL Database
- ChromaDB, Pypdf

- Built an intelligent document processing pipeline to extract key-value pairs and text from diverse formats using **Azure Document Intelligence** (Form Recognizer)
- Stored outputs in **SQL** & **vector databases** for semantic search, enabling downstream AI workflows.
- Integrated processed data into a **RAG application** for context-aware query answering
- Developed **Flask APIs** and optimized the pipeline for scalability, security, and high-volume document handling.
- Automated end-to-end document ingestion and retrieval workflow, improving processing efficiency and reducing manual effort in enterprise use cases.

Project: Document Extraction Agent & Automating with multi-Agentic system (June-25 to Sep-25)

Tools used:

- LangGraph
- Azure, Cromadb
- Flask, SQL

- Developed a **document extraction agent system** using LangChain and LangGraph, enabling to processing of input files, SOPs, and prompts to extract structured and unstructured data.
- Stored and managed extracted data in **SQL and Chroma vector databases**, integrating with a **RAG-based retrieval system** for context-aware query answering and intelligent information access.
- Finally, this is integrated with a multi-agent AI system for dynamic use across various applications.

Project: Heart Arrhythmia detection by using MIT-BIH ECG Bio signal (Sept-21 to Aprl-22)

Algorithms used:

- Python
- LSTM-CNN
- MIT-BIH data
- Deep learning

- Developed an arrhythmia detection system utilizing the MIT-BIH ECG bio signal dataset to identify and prevent heart-related and cardiovascular diseases.
- Implemented advanced signal processing and machine learning techniques to accurately detect various arrhythmias, enhancing early diagnosis and intervention.
- Conducted comprehensive validation and testing to ensure the system's reliability and accuracy, contributing to improved patient outcomes and preventive healthcare measures.

➢ **Additionally, I have worked on several other projects during my tenure that are not listed here for brevity.**

Declaration:

I hereby declare that all the information given above is true and correct to the best of my knowledge.

Date: Thanks & Regards,

Place: (M Penchala Prasad Rao)